



Machine Learning project

Multilayer Perceptron

Summary: This project is an introduction to artificial neural networks, with the implementation of a multilayer perceptron.

Version: 5.1

Contents

I	Introduction	2
I.1	A bit of history	2
I.2	Multilayer perceptron	3
I.3	Perceptron	4
II	Objectives	5
III	General instructions	6
IV	Mandatory part	7
IV.1	Foreword	7
IV.2	Dataset	8
IV.3	Implementation	9
IV.4	Submission	10
V	Bonus part	11
VI	Submission and Peer-Evaluation	12

Chapter I

Introduction

In the language of your choice, you are going to implement a **multilayer perceptron** in order to predict whether a cancer is malignant or benign based on a dataset of breast cancer diagnoses in Wisconsin.

I.1 A bit of history

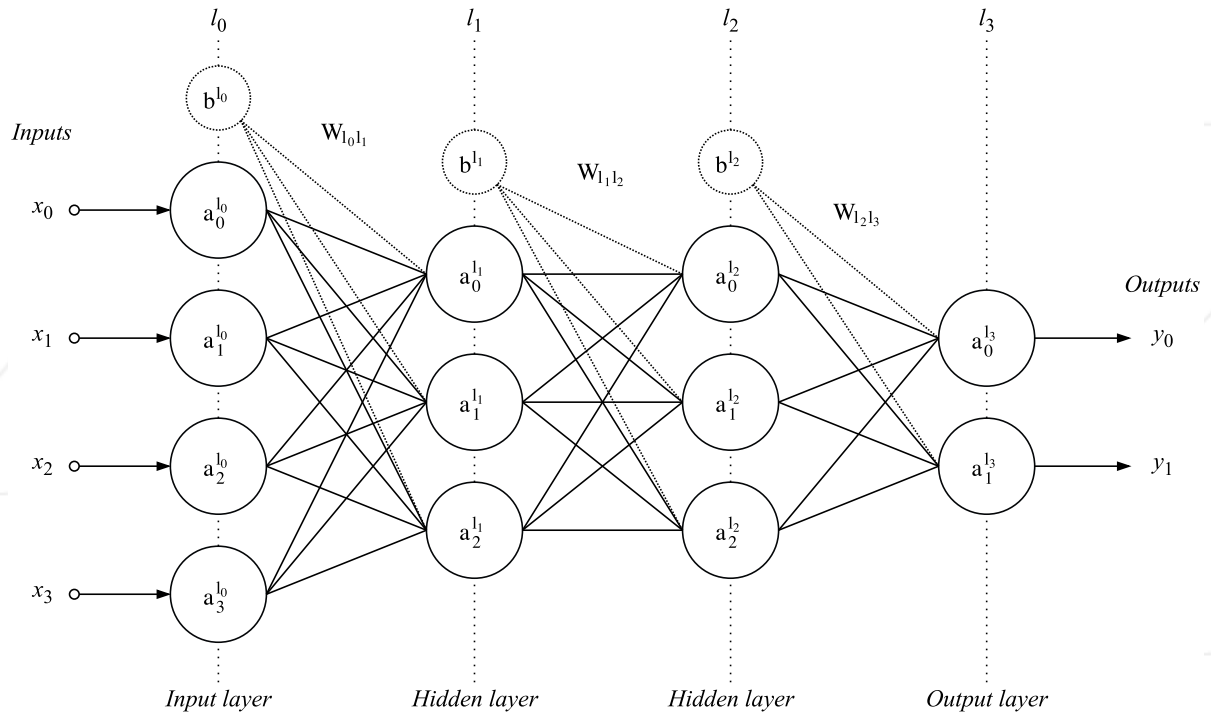
Machine learning is a vast field in which artificial neural networks are only a small subset. Nevertheless, we are going to tackle it since it is a really powerful tool that resurfaced a few years ago.

Contrary to what one may think, artificial neural networks have existed for a long time. In his 1948 paper '[Intelligent Machinery](#)', Alan Turing introduced a type of neural network called **B-type unorganised machine**, which he considered the simplest possible model of the nervous system.

The perceptron was invented by Frank Rosenblatt in 1957. It is a single-layer linear classifier and also one of the first neural networks to be implemented. Unfortunately, the results were not as good as expected, and the idea was abandoned. A bit more than 10 years later, the algorithm was improved as the **multilayer perceptron** and was used once again.

I.2 Multilayer perceptron

The multilayer perceptron is a feedforward network (meaning that the data flows from the input layer to the output layer) defined by the presence of one or more hidden layers, as well as an interconnection of all the neurons of one layer to the next.




The diagram above represents a network containing 4 dense layers (also called fully connected layers). Its inputs consist of 4 neurons and its output consists of 2 (perfect for binary classification). The weights of one layer to the next are represented by two-dimensional matrices noted $W_{l_j l_{j+1}}$. The matrix $W_{l_0 l_1}$ is of size $(3, 4)$ for example, as it contains the weights of the connections between layer l_0 and layer l_1 .

The bias is often represented as a special neuron which has no inputs and an output always equal to 1. Like a perceptron, it is connected to all the neurons of the following layer (the bias neurons are noted b^{l_j} on the diagram above). The bias is generally useful as it allows one to “control the behavior” of a layer.

I.3 Perceptron

The perceptron is the type of neuron that the multilayer perceptron is composed of. It is defined by the presence of one or more input connections, an activation function, and a single output. Each connection contains a weight (also called a parameter) which is learned during the training phase.



Two steps are necessary to get the output of a neuron. The first one consists in computing the weighted sum of the outputs of the previous layer with the weights of the input connections of the neuron, which gives

$$weighted\ sum = \sum_{k=0}^{N-1} (x_k \cdot w_k) + bias$$

The second step consists in applying an activation function on this weighted sum. The output of this function is the output of the perceptron and can be understood as the threshold above which the neuron is activated (activation functions can take many forms, and you are free to choose whichever one you want depending on the model to train. Here are some of the most frequently used ones to give you an idea: sigmoid, hyperbolic tangent, and rectified linear unit).

Chapter II

Objectives

The goal of this project is to give you a first approach to artificial neural networks, and to have you implement the algorithms at the heart of the training process. At the same time, you are going to have to get acquainted with the manipulation of derivatives and linear algebra, as they are indispensable mathematical tools for the success of the project.

Chapter III

General instructions

- This project will only be evaluated by humans. You are free to organize and name your files as you desire while respecting the restrictions listed below.
- You are free to use whatever language you want; you have no restrictions on that point.
- No libraries handling the implementation of artificial neural networks or the underlying algorithms are allowed; you must code everything from scratch. You can, however, use libraries to handle linear algebra and to display the learning curves.
- In the case of a compiled language, you must submit a Makefile. This Makefile must compile the project, and must contain the usual compilation rules. It should recompile and relink the program only as necessary. The dependencies should also be downloaded/installed with the Makefile as needed.
- The norm is not applied to this project. Nevertheless, you will be asked to be clear and structured in the conception of your source code.

Chapter IV

Mandatory part

IV.1 Foreword

A non-negligible part of the evaluation will be based on your understanding of the training phase (also called the learning phase) and the underlying algorithms. You will be asked to explain to your corrector the notions of **feedforward**, **backpropagation**, and **gradient descent**. Points will be awarded depending on the clarity of your explanations. These notions are important for the next projects in the branch and will represent a real asset if you wish to continue in this field.

IV.2 Dataset

The dataset is provided in the resources. It is a `csv` file with 32 columns, the `diagnosis` column being the label you want to predict, given all the other features of an example. It can be either the value M or B (for malignant or benign).

The features of the dataset describe the characteristics of a cell nucleus of breast mass extracted with [fine-needle aspiration](#). (For more detailed information, go [here](#)).

As you will see, there is an important data understanding phase before starting to implement the algorithm that will be able to classify it. A good practice would be to begin by exploring the dataset, displaying it with graphs, visualizing, and manipulating its different features.

You have to separate your dataset into two parts yourself: one for training and one for validation.



The data is raw and should be preprocessed before being used for the training phase.

IV.3 Implementation

Your neural network implementation must contain at least two hidden layers by default.

The idea is to make you write a program that is a bit more modular (you can use a file or directly provide arguments).

For example, you can add a hidden layer in the file:

```
cat example_network_modularity.txt

network = model.createNetwork([
    layers.DenseLayer(input_shape, activation='sigmoid'),
    layers.DenseLayer(24, activation='sigmoid', weights_initializer='heUniform'),
    layers.DenseLayer(24, activation='sigmoid', weights_initializer='heUniform'),
    layers.DenseLayer(24, activation='sigmoid', weights_initializer='heUniform'),
    layers.DenseLayer(output_shape, activation='softmax', weights_initializer='heUniform')
])

model.fit(network, data_train, data_valid, loss='categoricalCrossentropy', learning_rate=0.0314,
          batch_size=8, epochs=84)
```

Or for example with arguments:

```
python train.py --layer 24 24 24 --epochs 84 --loss categoricalCrossentropy --batch_size 8 --learning_rate
0.0314
```

You must also implement the **softmax** function on the output layer in order to obtain the output as a probabilistic distribution.

In order to evaluate the performance of your model in a robust way during training, you will split your dataset into two parts: one for training and one for validation (the validation dataset is used to determine the accuracy of your model on unknown examples).

You will also implement two learning curve graphs displayed at the end of the training phase (you are free to use any library you want for this purpose).

For example:




Figure IV.1: Loss

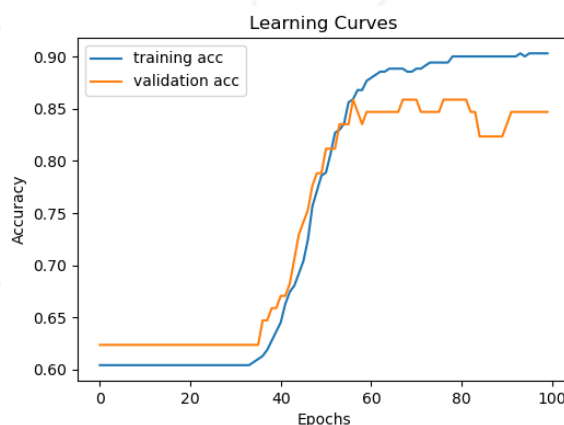


Figure IV.2: Accuracy

IV.4 Submission

You will submit three programs:

- A program to separate the dataset into two parts, one for training and the other for validation
- A training program
- A prediction program

(or you can submit a single program with an option to switch between the three phases)

To visualize your model's performance during training, you will display the training and validation metrics at each epoch.

For example:

```
python mlp.py --dataset data_training.csv
x_train shape : (342, 30)
x_valid shape : (85, 30)
epoch 01/70 - loss: 0.6882 - val_loss: 0.6788
...
epoch 39/70 - loss: 0.0750 - val_loss: 0.0406
epoch 40/70 - loss: 0.0749 - val_loss: 0.0404
epoch 41/70 - loss: 0.0747 - val_loss: 0.0400
...
epoch 70/70 - loss: 0.0640 - val_loss: 0.0474
> saving model './saved_model.npy' to disk...
```

- For the separate program, you are allowed to use a seed to obtain a repeatable result, because many random factors come into play (such as the weights and bias initialization).
- The training program will use **backpropagation** and **gradient descent** to learn on the training dataset and will save the model (network topology and weights) at the end of its execution.
- The prediction program will load the weights learned in the previous phase, perform a prediction on a given set (which will also be loaded), then evaluate it using the [binary cross-entropy error function](#) :

$$E = -\frac{1}{N} \sum_{n=1}^N [y_n \log p_n + (1 - y_n) \log(1 - p_n)]$$

Chapter V

Bonus part

The bonus part will be evaluated only if the mandatory part is perfectly done. You are free to implement any functionalities that you think could be interesting. Nevertheless, here is a non-exhaustive list of bonuses:

- A more complex optimization function (for example: Nesterov momentum, RMSprop, Adam, ...).
- A display of multiple learning curves on the same graph (really useful to compare different models).
- A history of the metrics obtained during training.
- The implementation of [early stopping](#).
- Evaluate the learning phase with multiple metrics.



The bonus part will only be assessed if the mandatory part is PERFECT. Perfect means the mandatory part has been completed in its entirety and works without malfunctioning. If you have not passed ALL the mandatory requirements, your bonus part will not be evaluated at all.

Chapter VI

Submission and Peer-Evaluation

Turn in your assignment in your `Git` repository as usual. Only the work inside your repository will be evaluated during the defense. Don't hesitate to double-check the names of your folders and files to ensure they are correct.